

Q3 2026 Financial & Technical Performance Report

1. Executive Summary

During the third quarter of 2026, our enterprise AI division successfully deployed an advanced Agentic Retrieval-Augmented Generation (RAG) platform. The platform incorporates multi-agent document classification, structural table extraction, and contextual retrieval algorithms.

Overall system throughput expanded by 45% compared to Q2, driven primarily by optimized vector indexing in PostgreSQL using the pgvector extension and HNSW indexing structures.

2. Technical Architecture & Ingestion Pipeline

The document ingestion pipeline processes incoming raw PDF files through four specialized agents. First, the Classifier Agent determines document domain types such as financial filings or technical papers. Second, the Chunker Agent separates tabular structures from narrative text to preserve table integrity.

Third, the Context Enrichment Agent generates situating context sentences for each chunk prior to embedding. Finally, a Critic Agent samples chunk output to evaluate coherence and retrieval fidelity.

3. Quarterly Performance Metrics Table

Quarter	Total Revenue (\$M)	RAG Latency (ms)	Active Users	Accuracy Rate
Q1 2026	\$12.4M	240 ms	14,500	94.2%
Q2 2026	\$15.8M	180 ms	22,100	96.5%
Q3 2026	\$21.3M	115 ms	38,900	98.7%

4. Conclusion and Future Strategic Investments

Looking forward to Q4, the engineering team plans to implement automatic retry loops in the Critic Agent and introduce LangGraph state graph orchestration for atomic manifest handoffs.